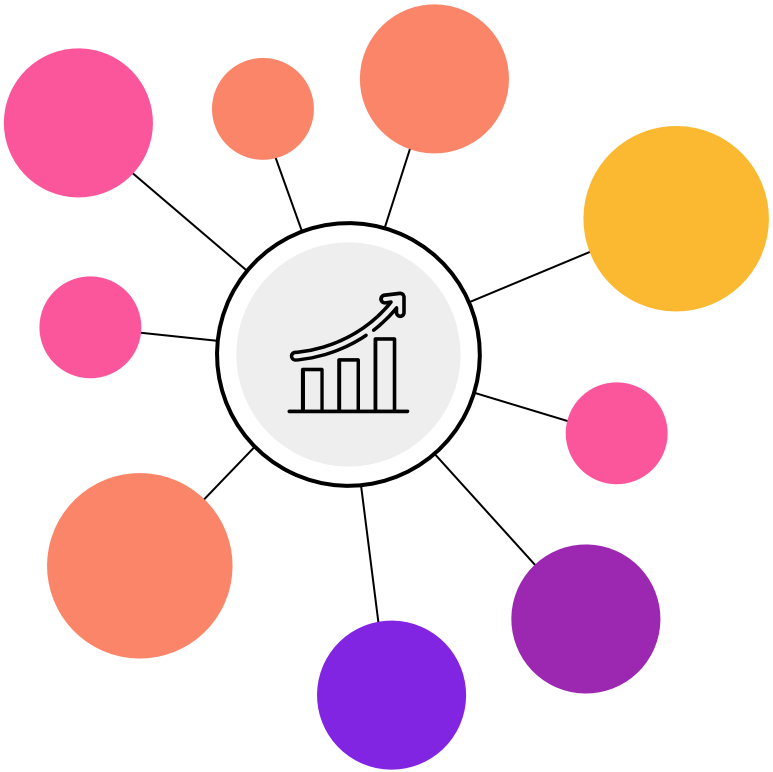

Software Engineer Salary

21KHDL1 - Group 7

Group Information

Members:
21127278 - Nguyen Trong Hieu
21127365 - Phan Phuoc Hai Nam
21127697 - Le Thi Minh Thu

Github: SWE-Salary (68 commits)

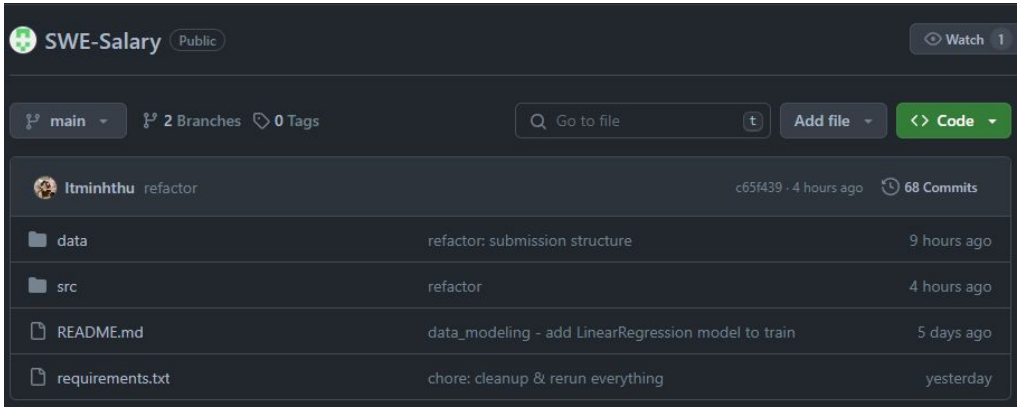

Software Engineer Salary

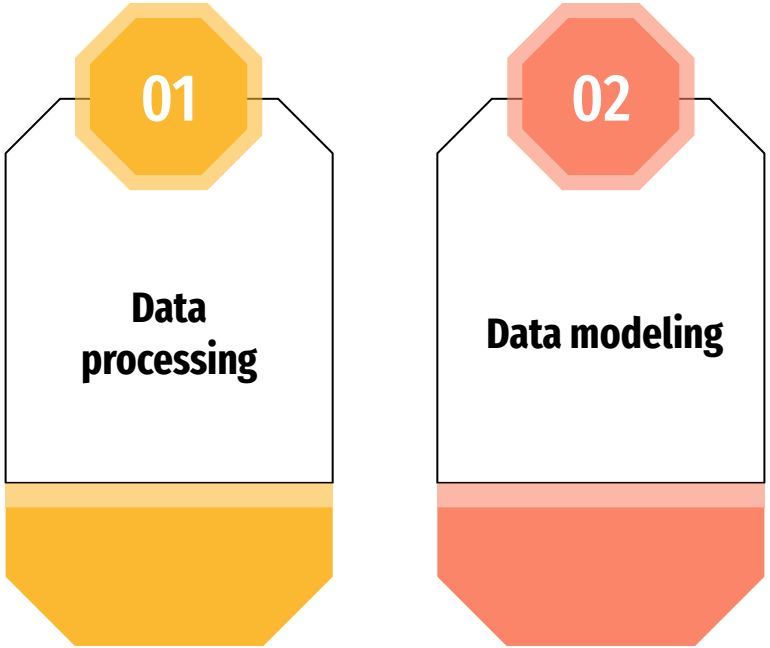

Software Engineer Salary

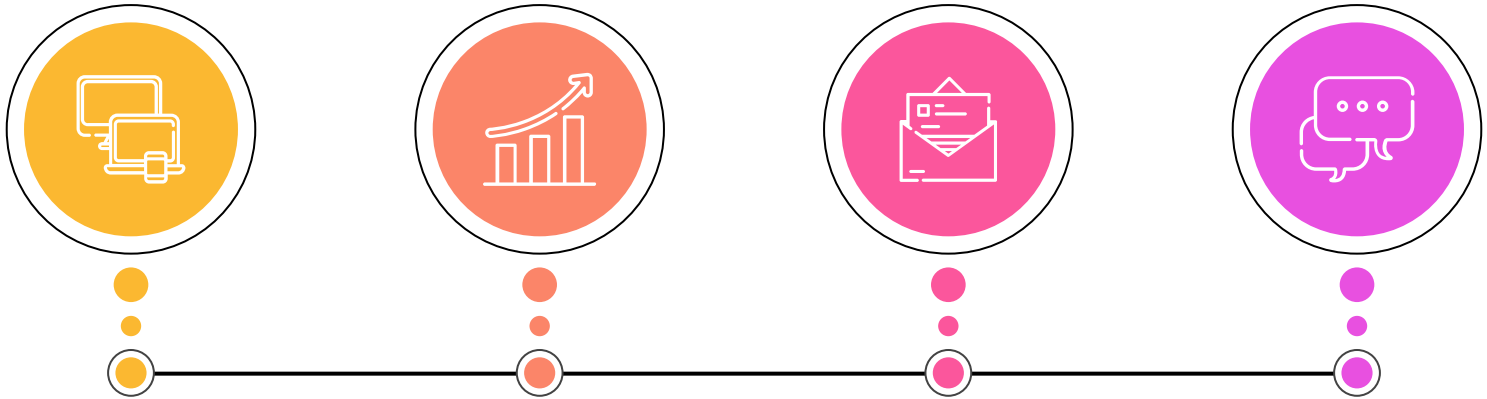

Collecting Data

Pre-processing

Exploration

Visualization

Data Collection

01

Define the needed features

02

Find resources (APIs, web or static data)

03

Locate the needed data

04

Get the data

Data Collection

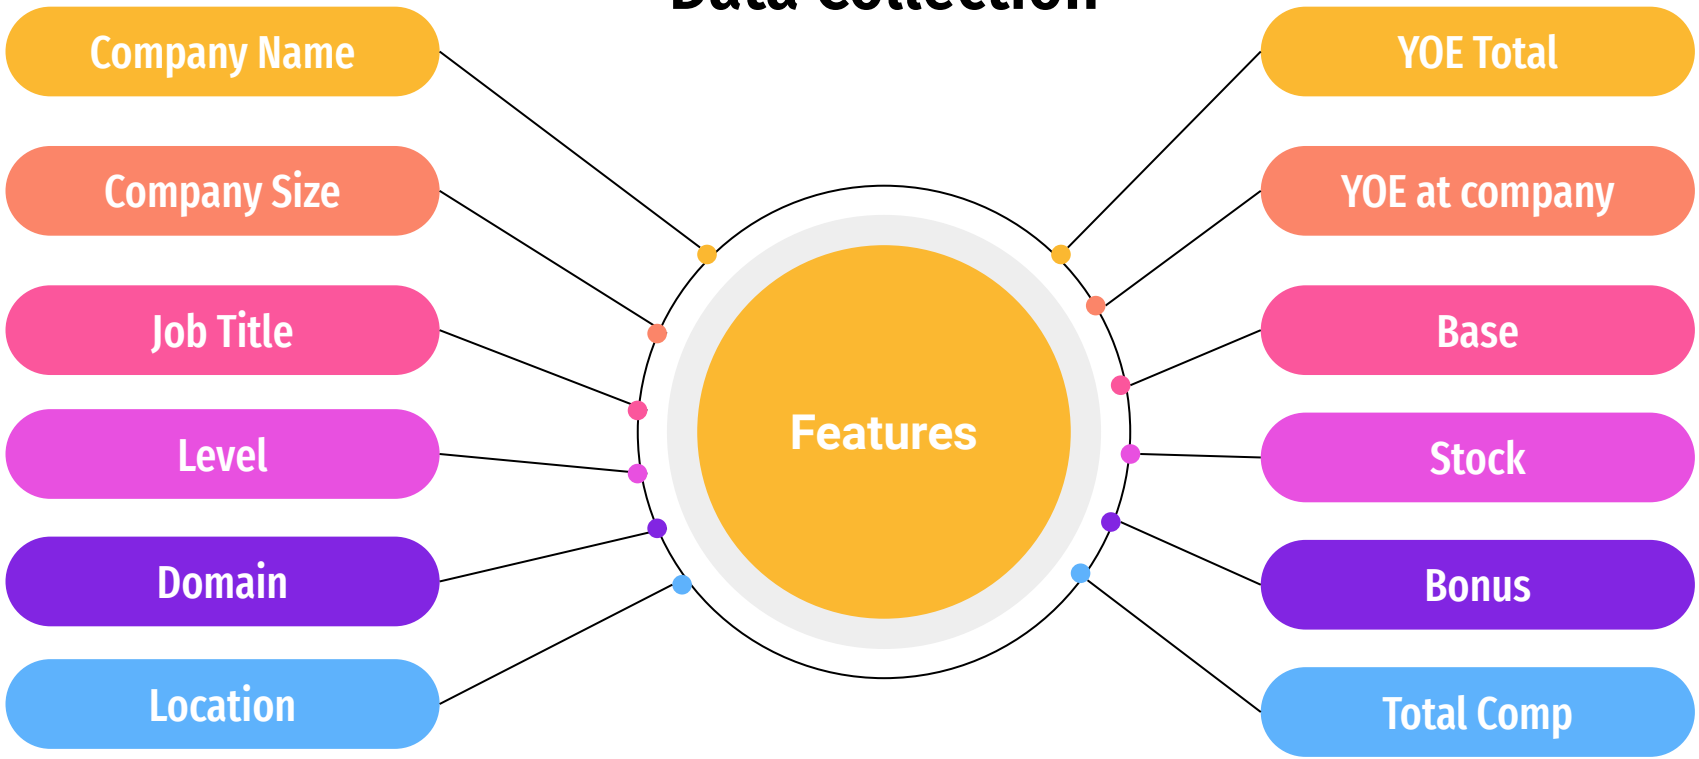

Data Collection

Resources

All the data is collected from levels.fyi:

- In 2022, levels.fyi had collected over 150.000 salary submissions.
- Their data comes from companies around the world.
- The information is well-defined and has a clear structure.

Data Collection

Resources

Company Location | Date	Level Name Tag	Years of Experience Total / At Company	Total Compensation (USD) Base | Stock (yr) | Bonus
Google San Francisco, CA | a day ago	L5 TPM	15 yrs 6 yrs	$350,000 200K | 120K | 30K
Google Mountain View, CA | 2 days ago	L5 AI	5 yrs 1 yr	$247,000 180K | 40K | 27K
Google Mountain View, CA | 5 days ago	L5 General	25 yrs 6 yrs	$300,000 185K | 85K | 30K
Google Boston, MA | 12/07/2023	L5 TPM	12 yrs 2 yrs	$321,000 171K | 120K | 30K
Google New York, NY | 12/05/2023	L5 Project Management	18 yrs 2 yrs	$323,000 194K | 100K | 29K

Data Collection

Locate Needed Data

CHALLENGE: All the needed data is not located in one site
→ We have to crawl data from different endpoints

`/companies/:id` → Company Name, Company Size

`/companies/:id/salaries` → Job Titles

`/companies/:id/salaries/:id` → Level, Domain, YOE, YOE at company, Base Salary, Stock, Bonus, Total Compensation, Location

Data Collection

Get Data

CHALLENGE: They have throttled the number of queries
→ We have to crawl data from different times

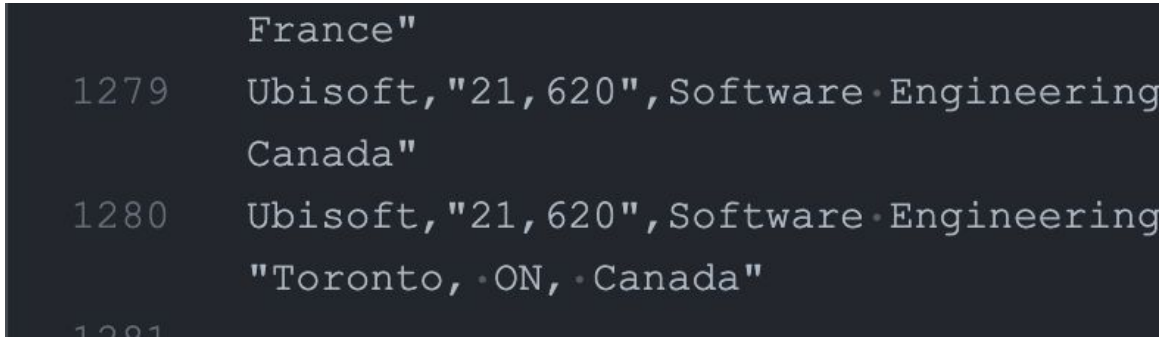

1279 records were crawled within 2 days because levels.fyi throttled each IP, allowing only approximately 50 access attempts every 30 minutes.

Data Preprocessing

01

Analyze the raw data

02

Handle missing data

03

Remove duplicates

04

Filter outliers

Data Preprocessing

Due to the well-structured source data and a small dataset of just over 1,000 records, our preprocessing will primarily involve data cleaning, eliminating the need for complex procedures like data integration or transformation.

One of the main tasks is to standardize the `YOE` and `YOE at company` into a format:

String type

`n yrs`
`n+ yrs`
`m-n yrs`

Int type

`n`

Data Preprocessing

```
Missing values ratio:

company                0.0
company_size           0.0
job_title              0.0
level                  0.0
domain                 0.0
yoe_total              0.0
yoe_at_company         0.0
base                   0.0
stock                 26.2
bonus                 44.7
total_compensation     0.0
location               0.0
dtype: float64
```

With a missing ratio of 26.2% for the stock column and 44.7% for the bonus column, despite the very high missing data rates, these are not data that every company can provide publicly.

Moreover, they are used to validate the `base` and `total_compensation`, so we have decided to fill the missing values with 0.

Data Preprocessing

After performing additional checks such as ensuring the validity of `total_compensation` and converting numerical data to integers and floats, we obtain the pre-processed data as follows.

```
Data columns (total 12 columns):
 #   Column              Non-Null Count  Dtype
---  ------              --------------  -----
 0   company             1250 non-null   object
 1   company_size        1250 non-null   int64
 2   job_title           1250 non-null   object
 3   level               1250 non-null   object
 4   domain              1250 non-null   object
 5   yoe_total           1250 non-null   int64
 6   yoe_at_company      1250 non-null   int64
 7   base                1250 non-null   float64
 8   stock               1250 non-null   float64
 9   bonus               1250 non-null   float64
 10  total_compensation  1250 non-null   float64
 11  location            1250 non-null   object
dtypes: float64(4), int64(3), object(5)
memory usage: 127.0+ KB
```

```
raw_df.apply(missing_ratio)
✓ 0.0s

company              0.0
company_size         0.0
job_title            0.0
level                0.0
domain               0.0
yoe_total            0.0
yoe_at_company       0.0
base                 0.0
stock                0.0
bonus                0.0
total_compensation   0.0
location             0.0
dtype: float64
```

Data Exploration

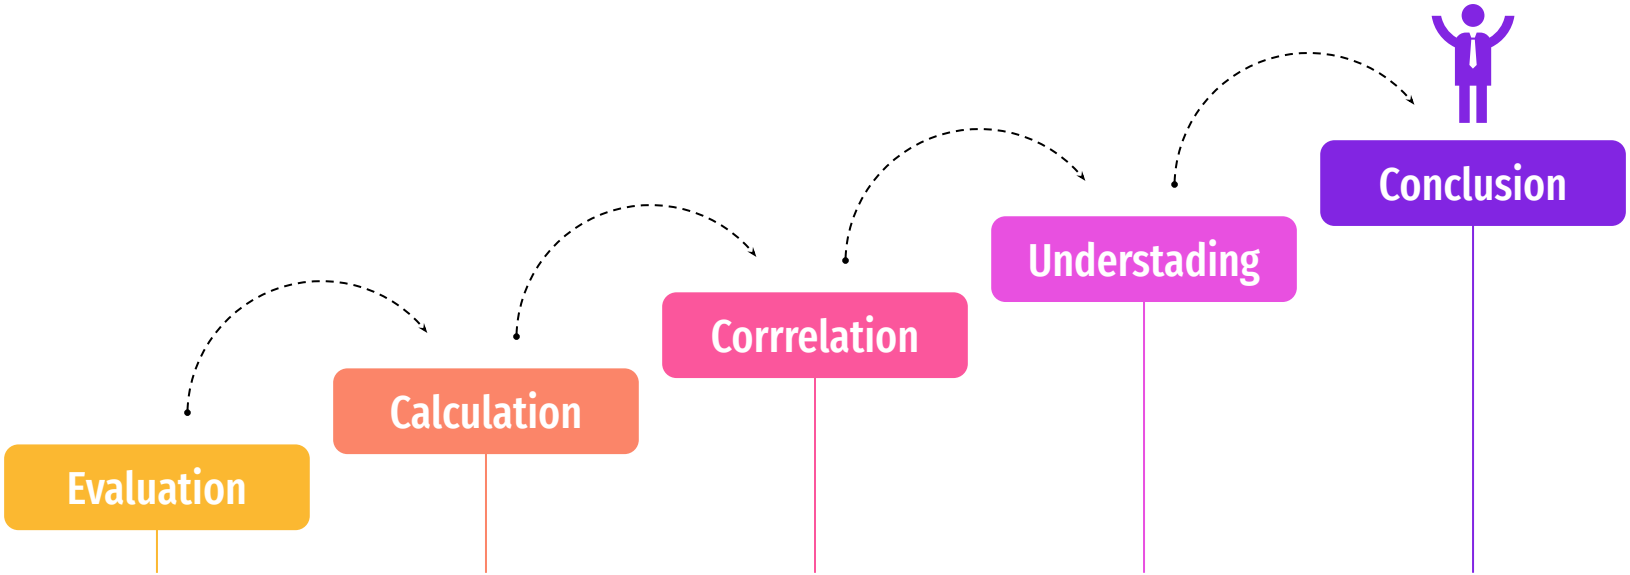

Evaluation

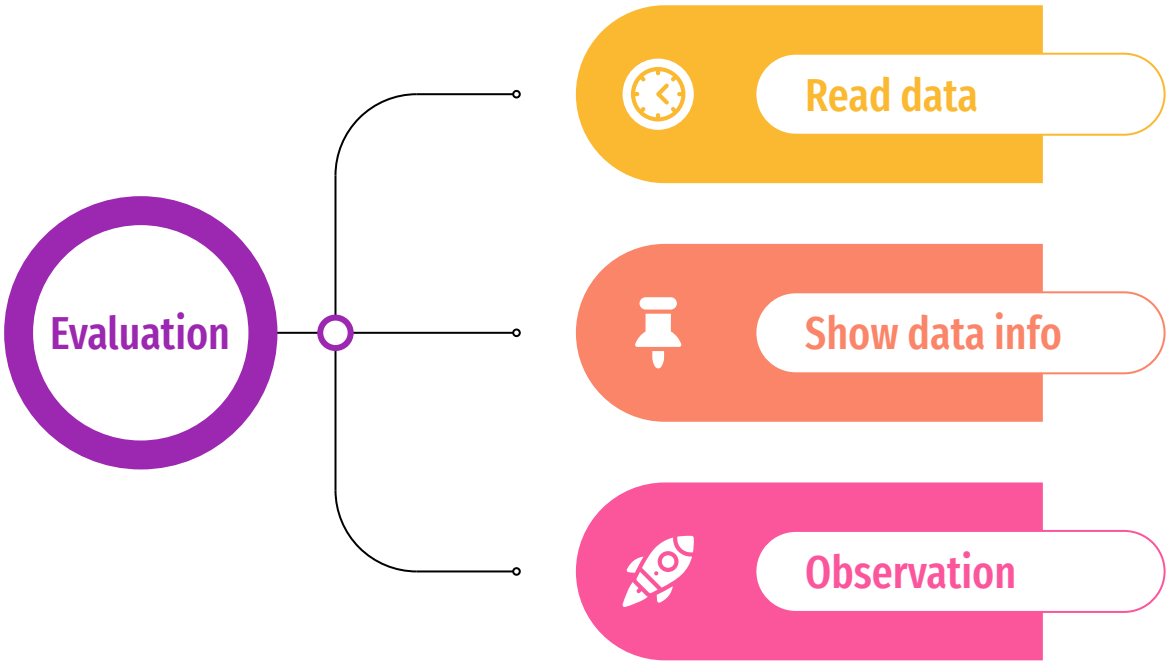

Read data from file .csv and print data

Print data info: name of attributes, non-null count and dtypes

Give observation

Evaluation

Read data

```
cleaned_data = pd.read_csv('../data/cleaned_data.csv')
cleaned_data
```

Python

	company	company_size	job_title	level	domain	yoe_total	yoe_at_company	base	stock	bonus	total_compensation	location
0	Logitech	7250	Software Engineer	I4	Testing (SDET)	10	5	190000.0	10000.0	0.0	200000.0	San Francisco Bay Area
1	Logitech	7250	Software Engineer	I2	ML / AI	4	3	126000.0	0.0	7000.0	133000.0	Vancouver, WA
2	Logitech	7250	Software Engineer	I3	Testing (SDET)	11	11	120000.0	5000.0	12000.0	137000.0	San Francisco, CA
3	Logitech	7250	Software Engineer	I4	Production	8	8	100000.0	10000.0	0.0	110000.0	Hsin-chu, TP, Taiwan
4	Logitech	7250	Software Engineer	I1	ML / AI	2	0	123100.0	0.0	0.0	123100.0	New York, NY
...	...	...	...	...	...	...	...	...	...	...	...	...
1245	Ubisoft	21620	Software Engineering Manager	L3	Video Game	13	13	80400.0	0.0	0.0	80400.0	Bordeaux, AQ, France
1246	Ubisoft	21620	Software Engineering Manager	L4	API Development (Back-End)	10	1	115300.0	0.0	11500.0	126800.0	Montreal, QC, Canada
1247	Ubisoft	21620	Software Engineering Manager	L3	Other	12	12	73900.0	0.0	0.0	73900.0	IL, France
1248	Ubisoft	21620	Software Engineering Manager	L4	Full Stack	20	9	125000.0	0.0	5000.0	130000.0	M
1249	Ubisoft	21620	Software Engineering Manager	L4	Game Development	25	10	100000.0	20000.0	2000.0	122000.0	To

1250 rows × 12 columns

Evaluation

Data info

```
cleaned_
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1250 entries, 0 to 1249
Data columns (total 12 columns):
 #   Column              Non-Null Count  Dtype
---  ------              --------------  -----
 0   company             1250 non-null   object
 1   company_size        1250 non-null   int64
 2   job_title           1250 non-null   object
 3   level               1250 non-null   object
 4   domain              1250 non-null   object
 5   yoe_total           1250 non-null   int64
 6   yoe_at_company      1250 non-null   int64
 7   base                1250 non-null   float64
 8   stock               1250 non-null   float64
 9   bonus               1250 non-null   float64
 10  total_compensation  1250 non-null   float64
 11  location            1250 non-null   object
dtypes: float64(4), int64(3), object(5)
memory usage: 117.3+ KB
```

Evaluation

Observation

Observation:

- The data has total 12 columns and 1250 rows
- The data has no missing values
- The total data size is higher than 1000 which means it a well collecting data
- The type of the data is float64 and int64 which means it is a numerical data so we can easily apply some statistical methods to explore and analyze the data

Calculation

Calculate numeric data

Calculate mean, median, lower quartile, upper quartile and mode of numeric data

Give observation

Give observation about mean, median, mode, max, min

Calculation

Calculate

```python
def mean(df):
    return (df.mean()).round(1)

def missing_ratio(s):
    return (s.isna().mean() * 100).round(1)

def median(df):
    return (df.quantile(0.5)).round(1)

def lower_quartile(df):
    return (df.quantile(0.25)).round(1)

def upper_quartile(df):
    return (df.quantile(0.75)).round(1)

def mode(df):
    return df.mode().iloc[0]

num_col_info_df = cleaned_data.select_dtypes(include=np.number)
num_col_info_df = num_col_info_df.agg([mode, mean, missing_ratio, "min", lower_quartile, median, upper_quartile, "max"])
num_col_info_df
```

	company_size	yoe_total	yoe_at_company	base	stock	bonus	total_compensation
mode	865406.0	10.0	2.0	200000.0	0.0	0.0	200000.0
mean	209016.8	9.5	3.5	163774.0	74274.5	16002.5	264438.1
missing_ratio	0.0	0.0	0.0	0.0	0.0	0.0	0.0
min	570.0	0.0	0.0	1100.0	0.0	0.0	6300.0
lower_quartile	19410.0	5.0	1.0	109925.0	0.0	0.0	134325.0
median	94520.0	9.0	2.0	157000.0	30000.0	7150.0	206750.0
upper_quartile	212570.0	13.0	4.0	200000.0	90000.0	25000.0	320875.0
max	865406.0	37.0	24.0	900000.0	750000.0	150000.0	2960000.0

Calculation

Observation

Observation

- Company size:
  - The minimumm company size is 570 which means all of the data collected from large companies
  - The maximum company size is 865,406 of Amazon which is the largest company in the world
  - The average company size is 309,147 which means the data mostly is collected from BIG-TECH companies in the world
  - The median company size is 147,000 which is much less than the average show that the top companies in the world have big difference in size compared to the orthers
- Years of experience in total:
  - The minimum years of experience is 0 which means there are some fresh graduated students can join in these large companies
  - The maximum years of experience is 37 which means this career is not only for youngster but also for the elder
  - The average years of experience is 11 which means the data mostly is collected from the people who have a lot of experience in this career and the experienced people are more likely to be hired by the top companies
  - The median years of experience is 8 which is a little bit less than the average show that most of people in big companies have a lot of experience in this career
- Years of experience in current company:
  - The minimum years of experience in current company is 0 which means there are some fresh graduated students can join in these large companies
  - The maximum years of experience in current company is 28 which means there are some people who have been working for a long time in the same company
  - The average years of experience in current company is 5 show that most of people in big companies stay there for a long time
  - The median years of experience in current company is 2 which is much less than the average
- Base salary:
  - The minimum
  - The maximum base salary is 900,000 which is a huge number even though this data is collected from the top companies in the world
  - The average base salary is 309,147 which is a huge number and it is not a surprise because the data is collected from the top companies in the world
  - The median base salary is 157,000 which is much less than the average show that the top companies in the world have big difference in salary compared to the orthers
- Stock:
  - The minimum is 0 which means there are some people who do not have stock or some companies do not offer stock to their employees
  - The maximum stock is 750,000 which is near the maximum base salary show that some companies offer a huge amount of stock to their employees instead of huge base salary
  - The average stock is 121,557 is so much less than the max base salary prove that the stock is not a huge part of the total compensation
  - The median stock is 0 which means most of companies do not offer stock to their employees
- Bonus:
  - The minimum bonus is 0 which means there are some companies do not offer bonus to their employees so that not all big companies offer a good treat to their employees
  - The maximum bonus is 275,000 which is nearly 4 months of maximum base salary
  - The average bonus is 43,521 which is much less than the max base salary show that the bonus is not a huge part of the total compensation
  - The median bonus is 0 which means most of companies do not offer bonus to their employees
- Total compensation:
  - The minimum
  - The maximum total compensation is 2,960,000 which is a super huge number compared to a business profit
  - The average total compensation is 538,607 which is much less than the max total compensation show that the top companies in the world have big difference in total compensation compared to the orthers
  - The median total compensation is 150,000 which is surprisely much less than the average show that the top companies in the world have big difference in total compensation compared to the orthers

Correlation

Show heatmap	Give observation
	
Create heatmap using numeric data to show correlation	Give observation about correlation of data attributes

Calculation

Heatmap

```
data_copy = cleaned_data.copy()

data_copy.drop(['company', 'job_title', 'level', 'domain', 'location'], inplace=True, axis='columns')

fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(data_copy.corr(), cmap='RdBu', center=0,ax=ax)
plt.show()
```

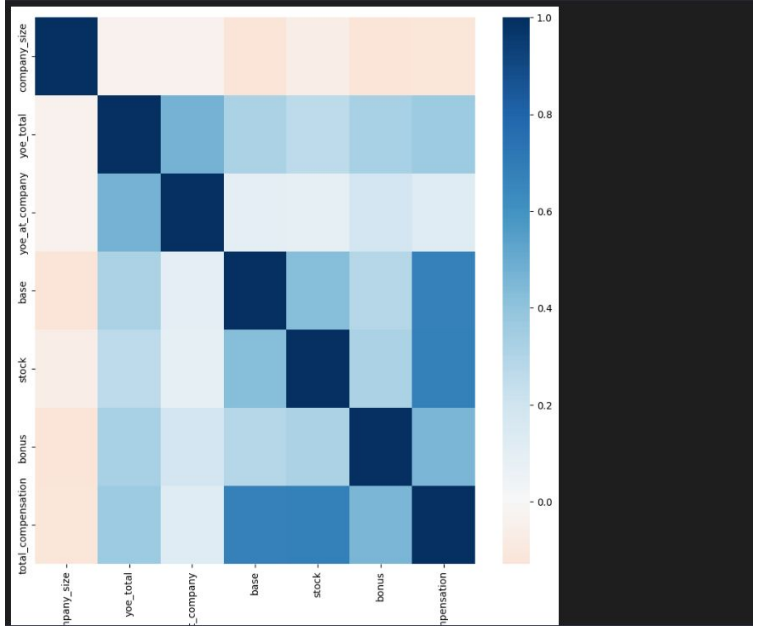

Calculation

Observation

Observation:

- The company size has a negative correlation with total compensation and base salary which means the bigger the company is, the less the total compensation and base salary are (especially the base salary)
- The years of experience has a positive correlation with total compensation and base salary which means the more the years of experience is, the more the total compensation and base salary are (especial[illegible] total compensation)
- The years of experience in current company has nearly no effect on total compensation and base salary

Understanding

Make questions	Answer questions
	
Make some questions about data	Answer questions to understand data

Understanding

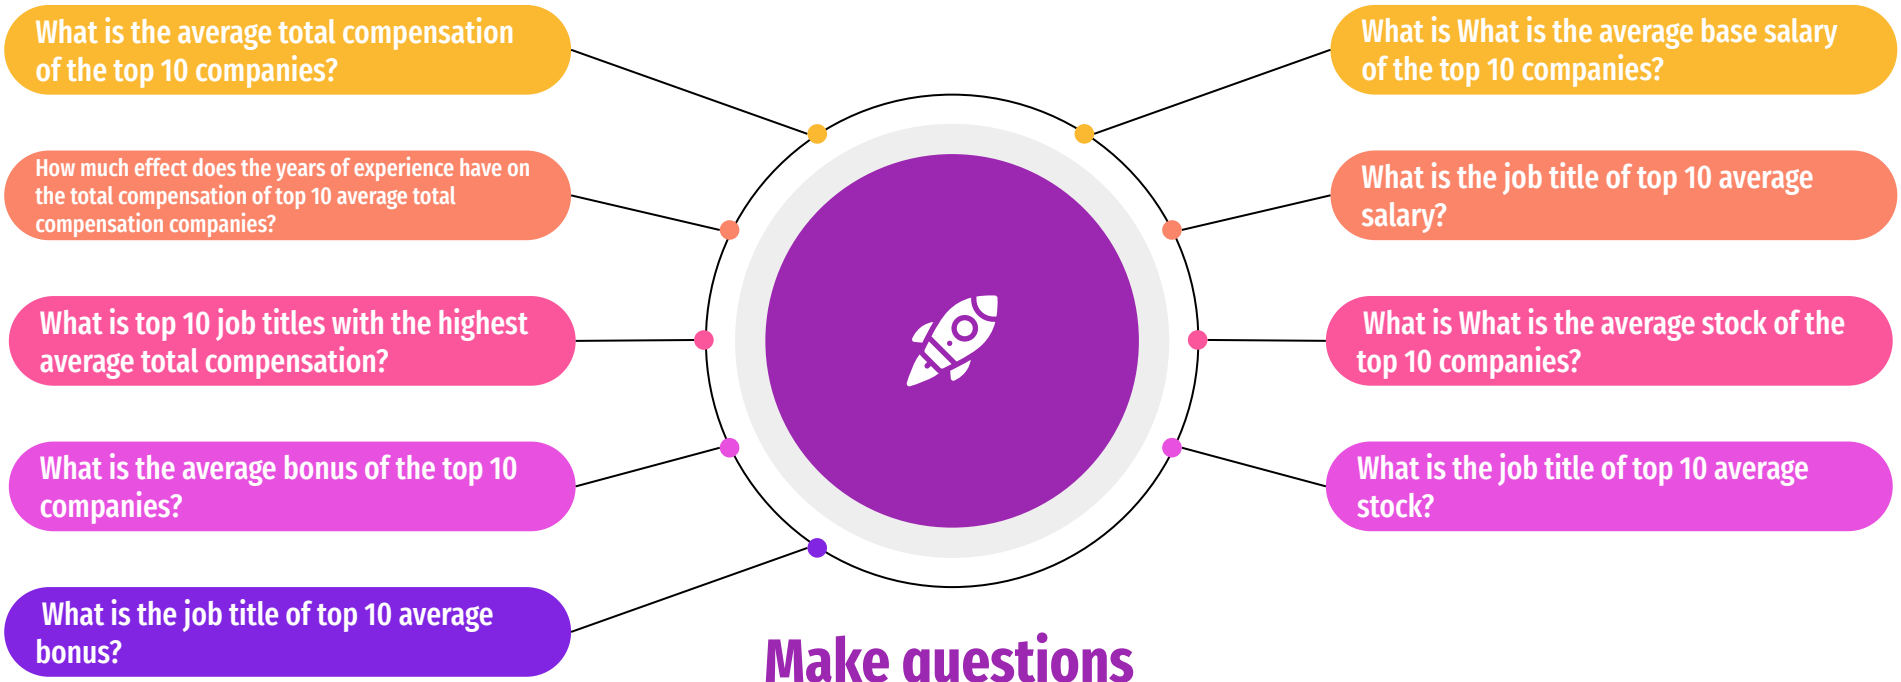

Make questions

Understanding

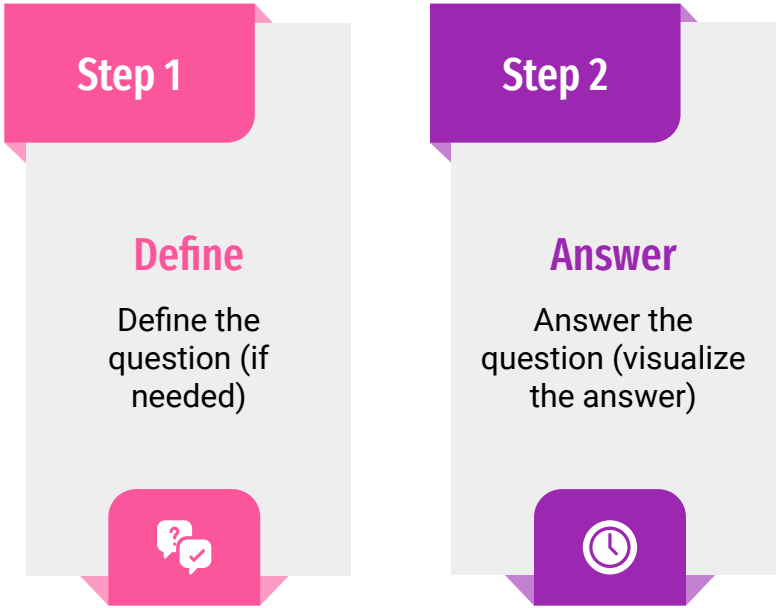

Answer questions

Understanding

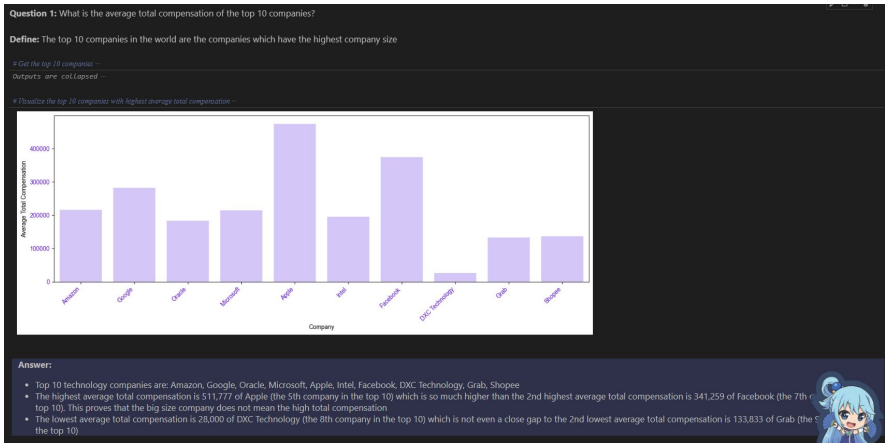

Answer questions

Visualization

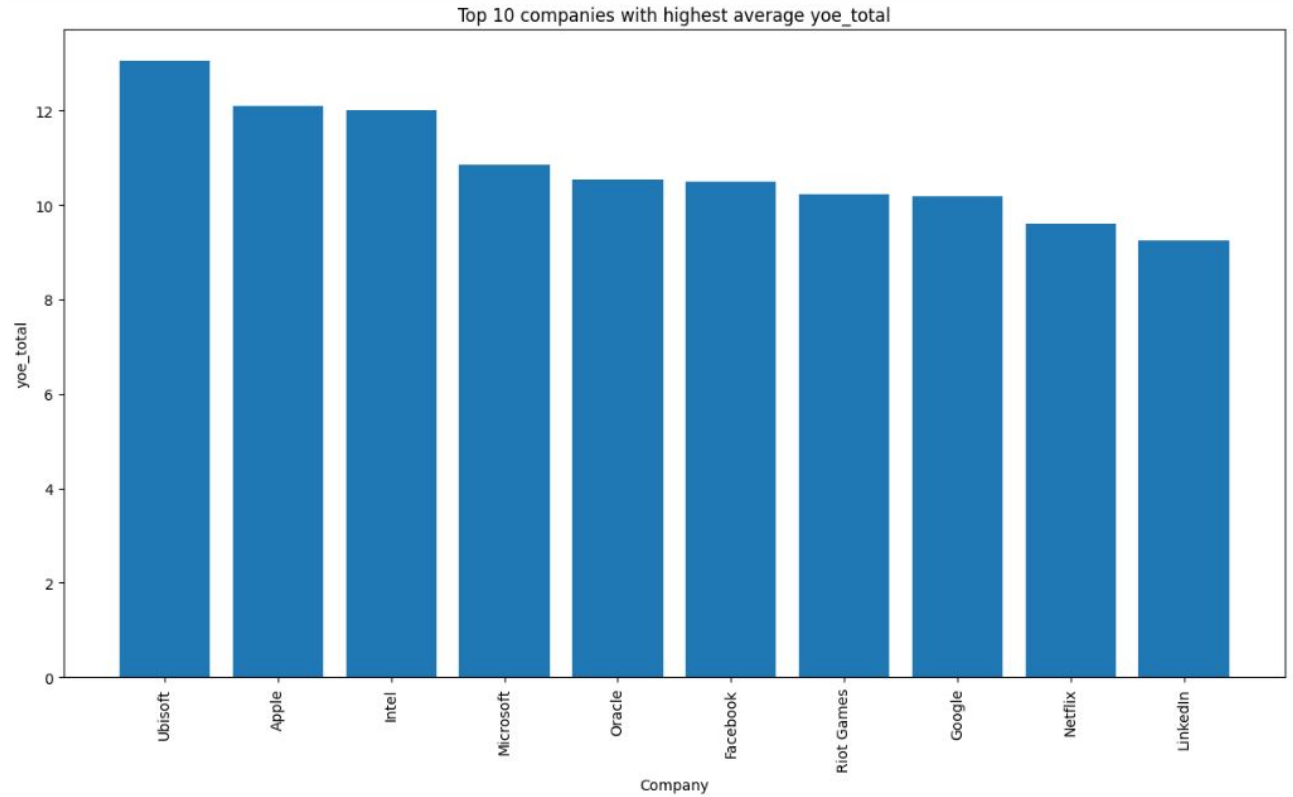

Visualization

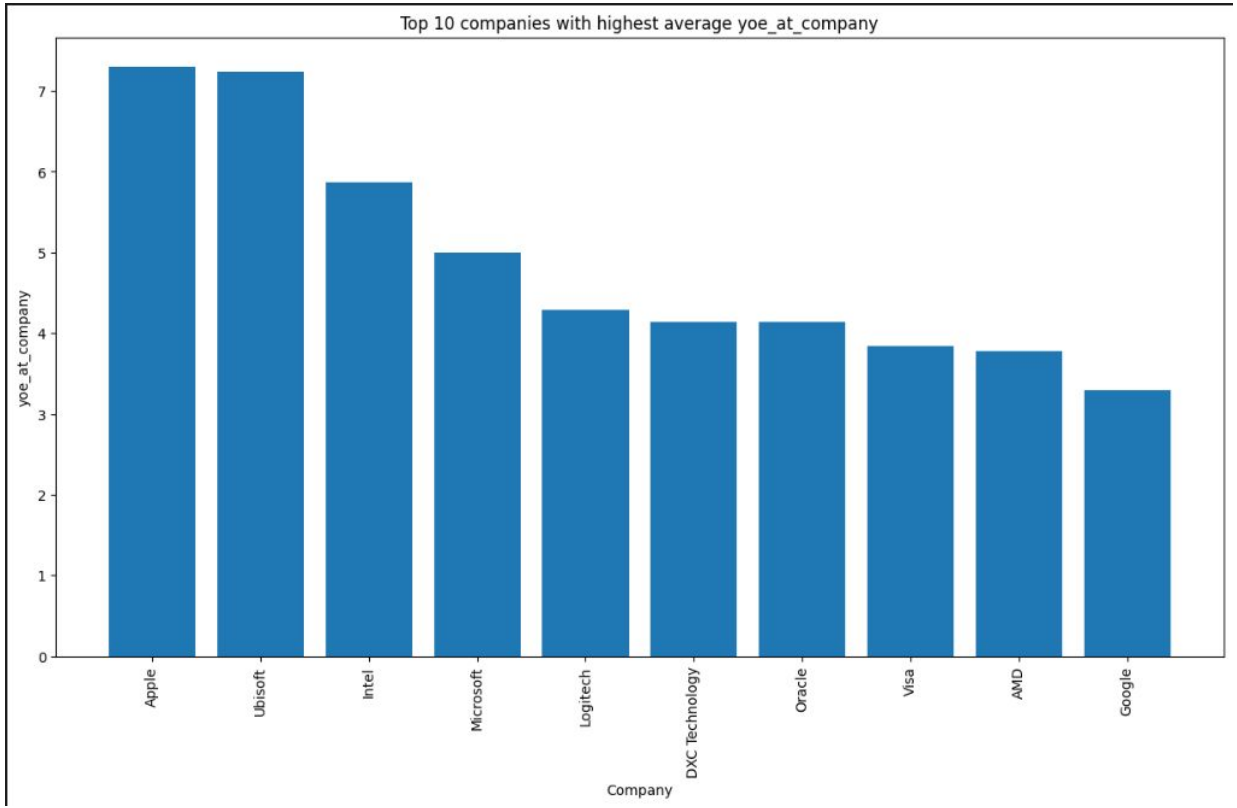

Visualization

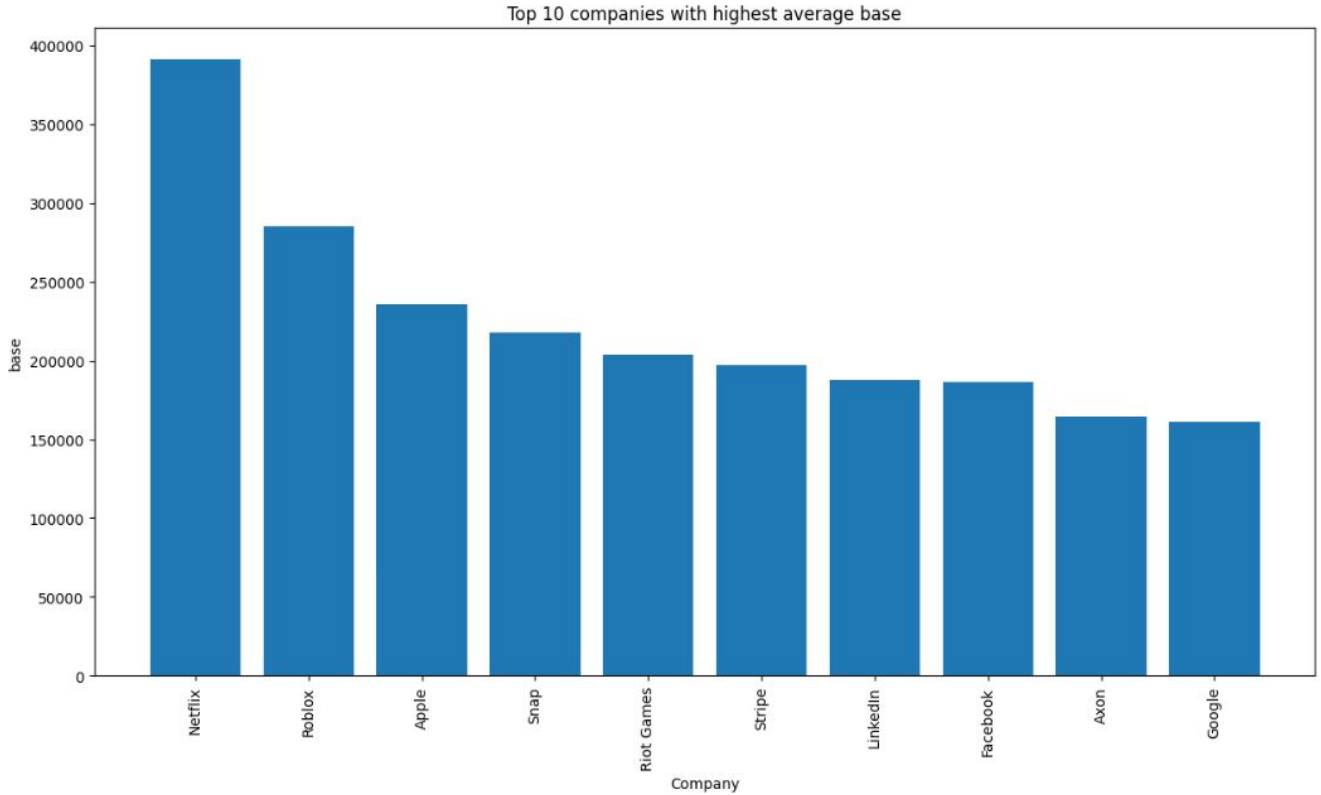

Visualization

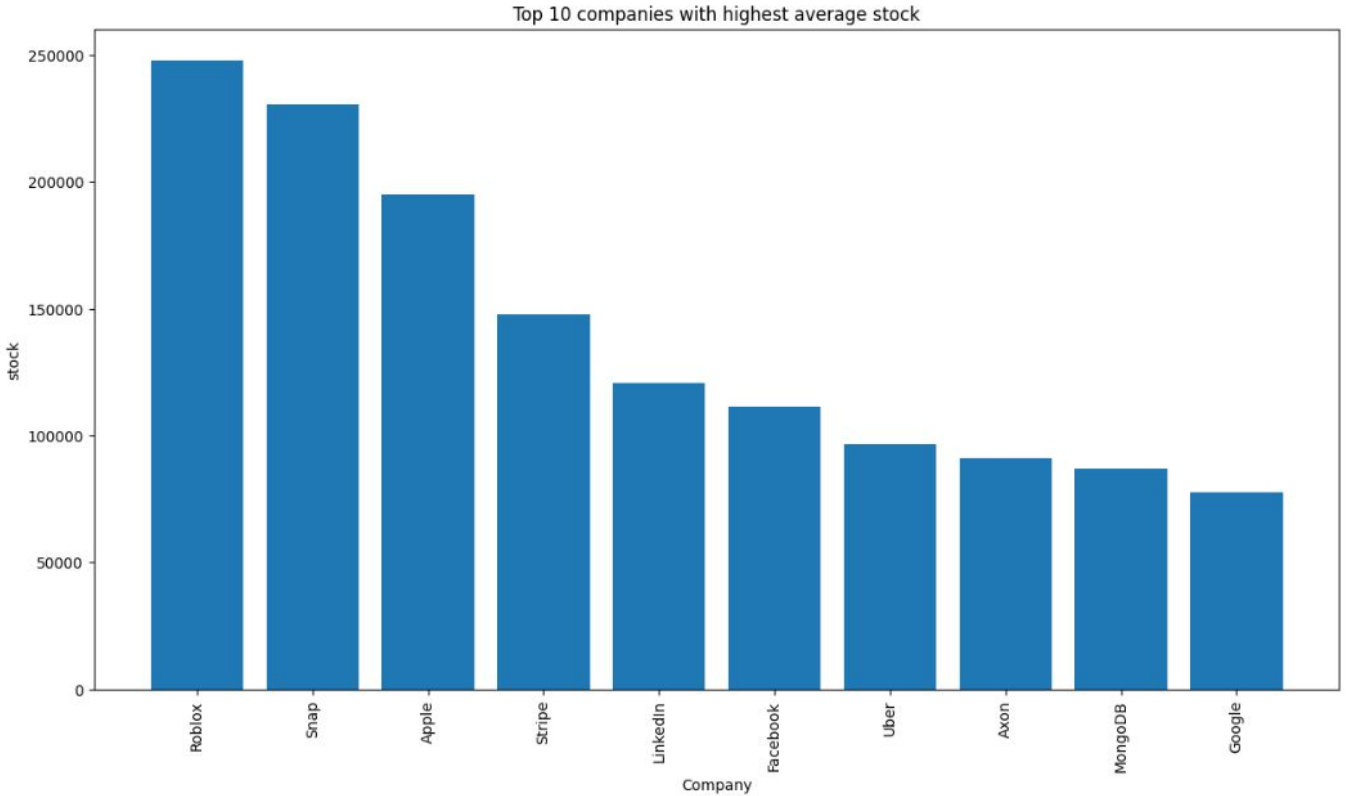

Visualization

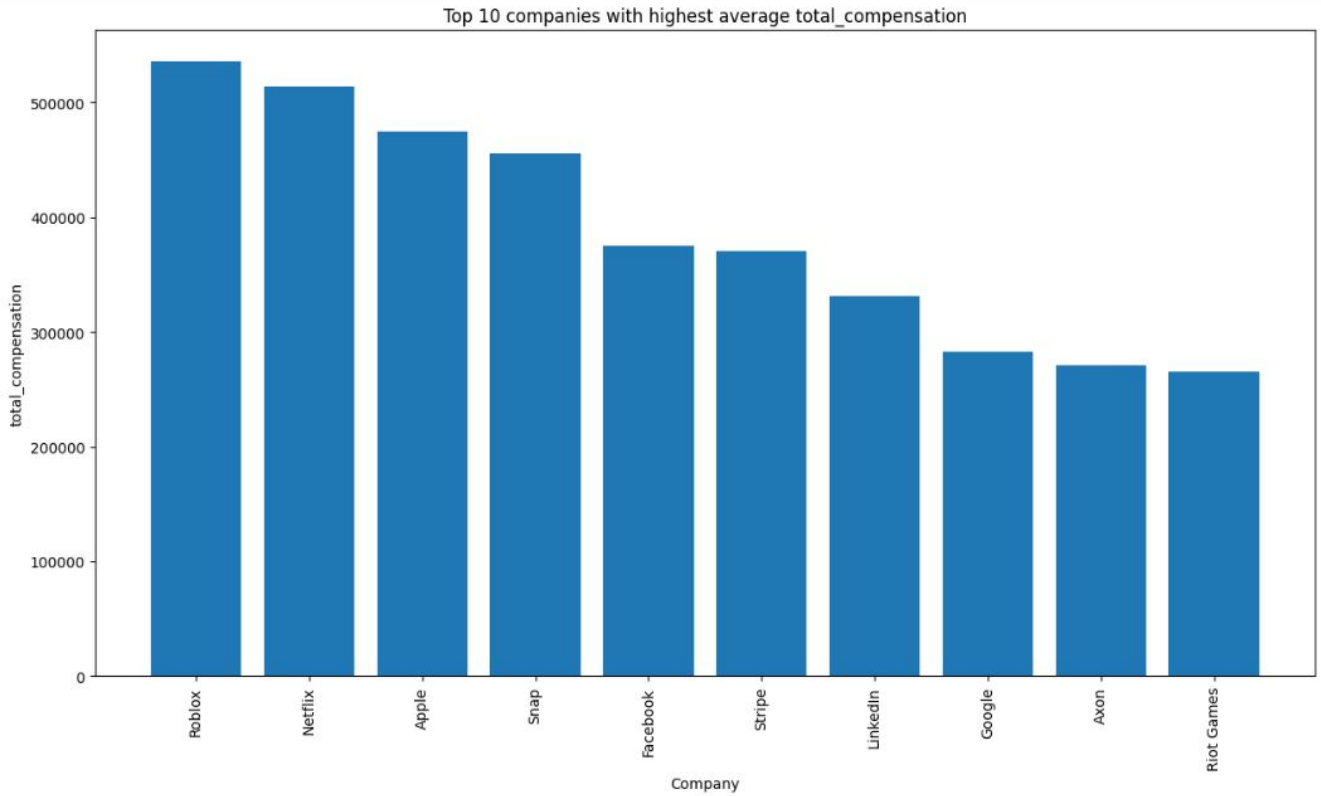

Visualization

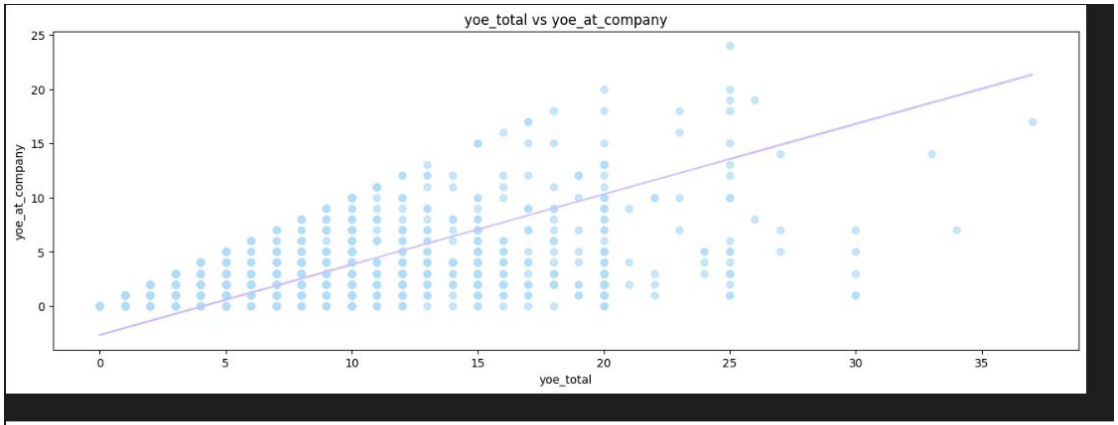

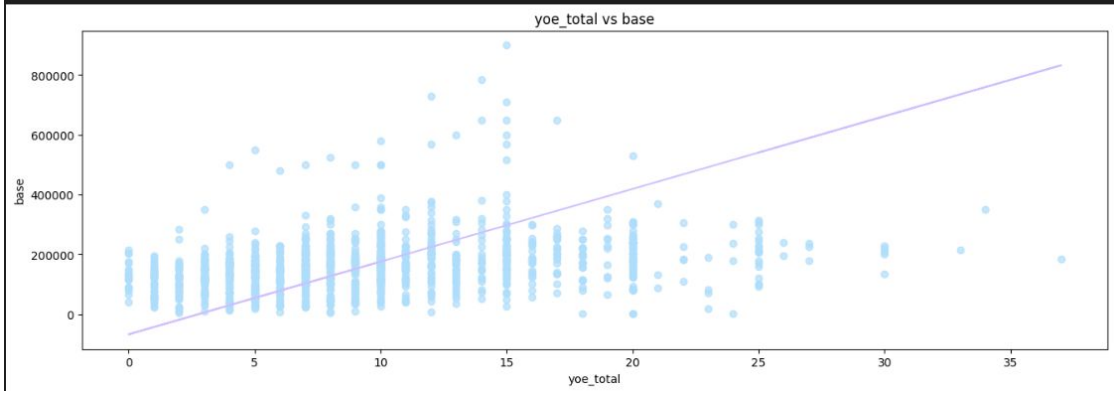

Visualization

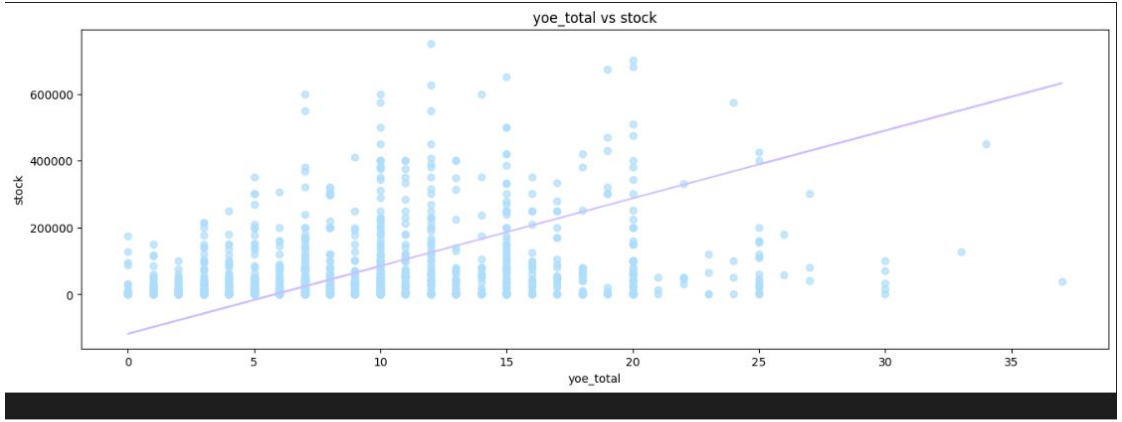

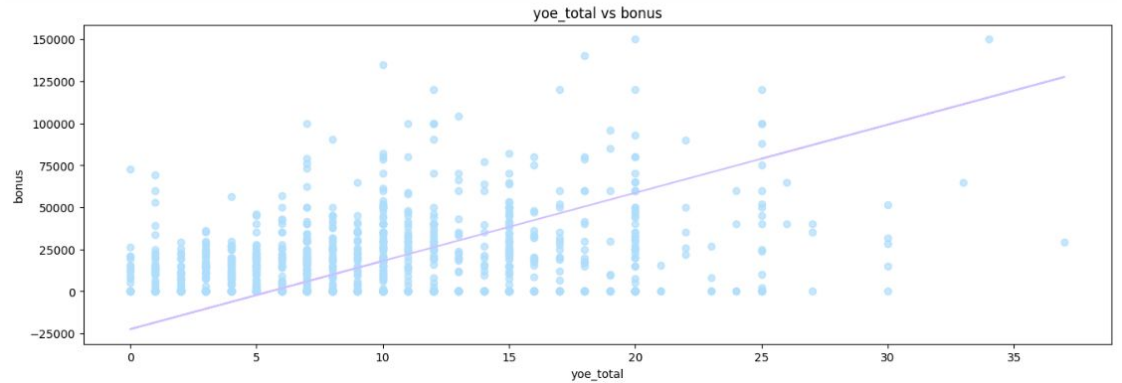

Visualization

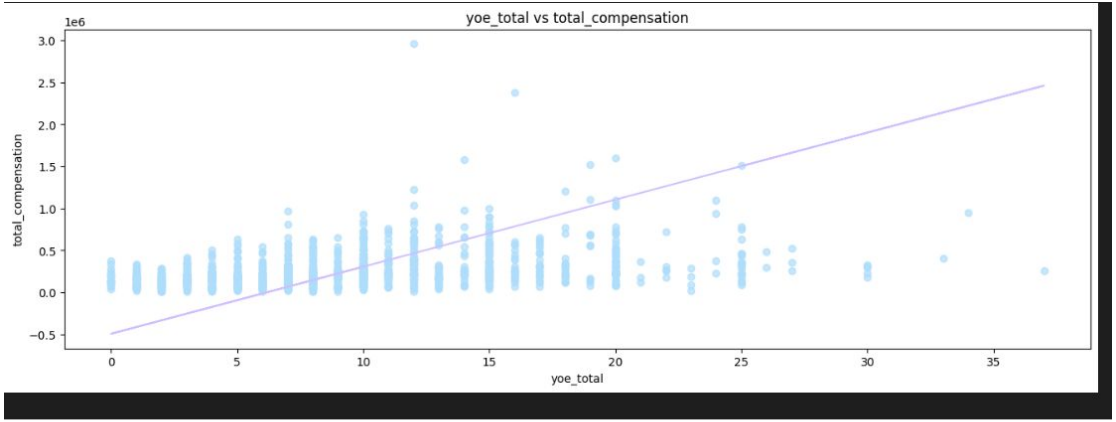

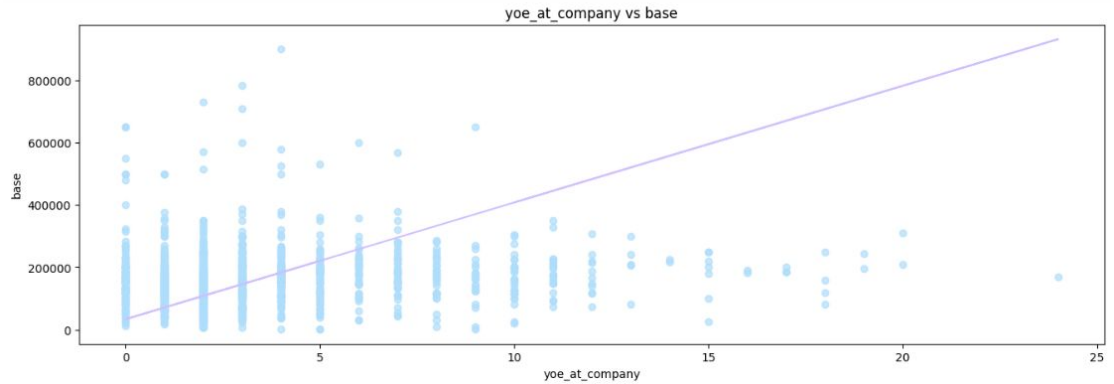

Visualization

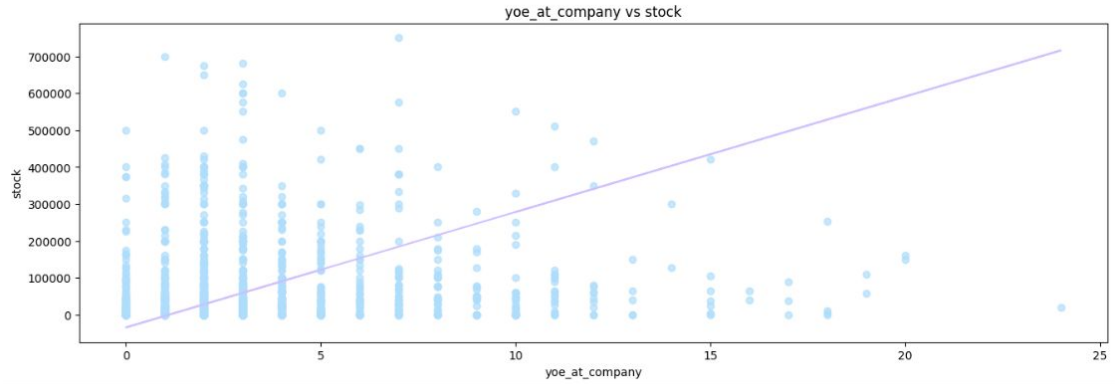

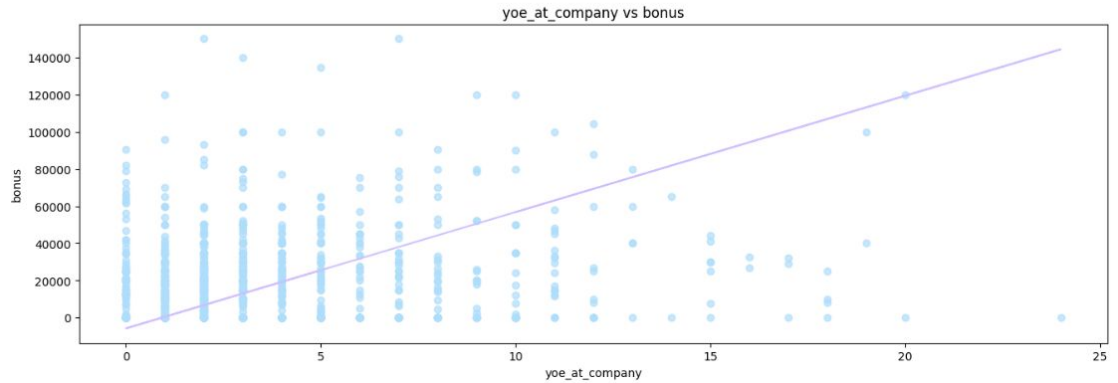

Visualization

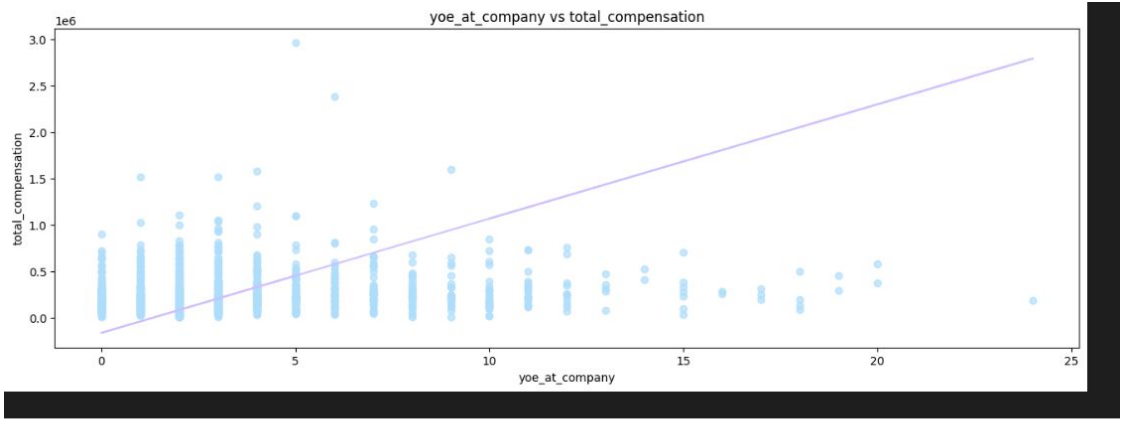

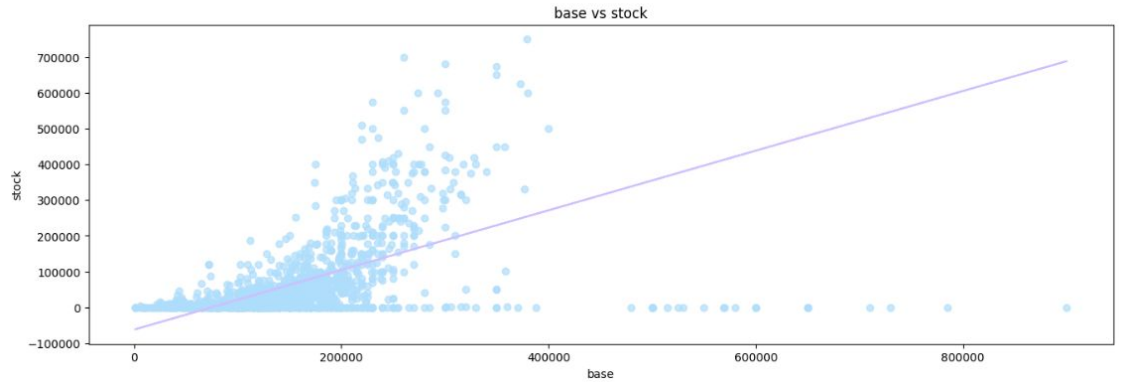

Visualization

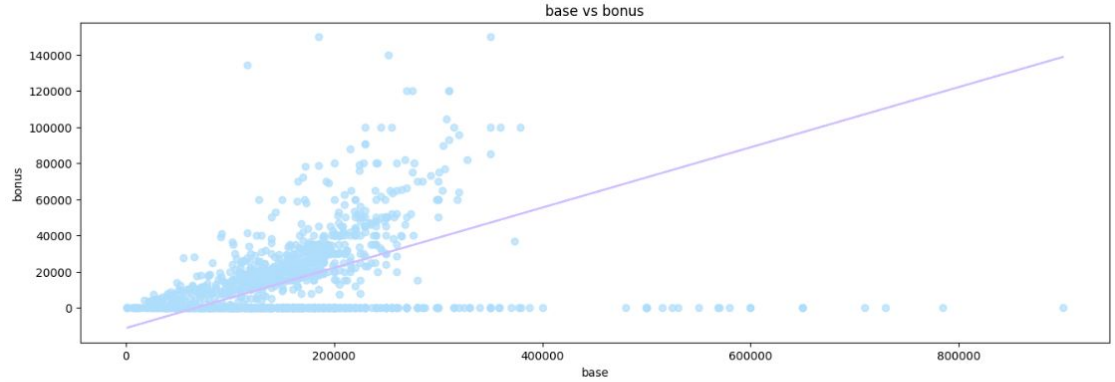

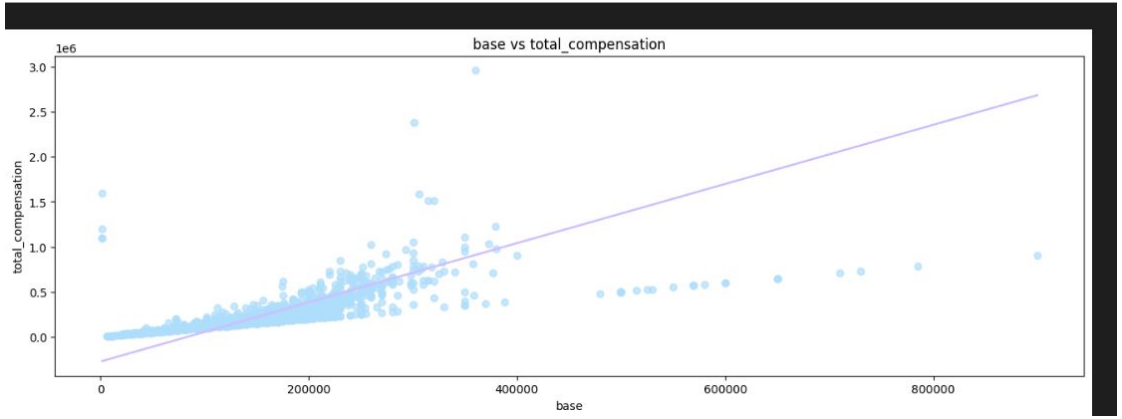

Visualization

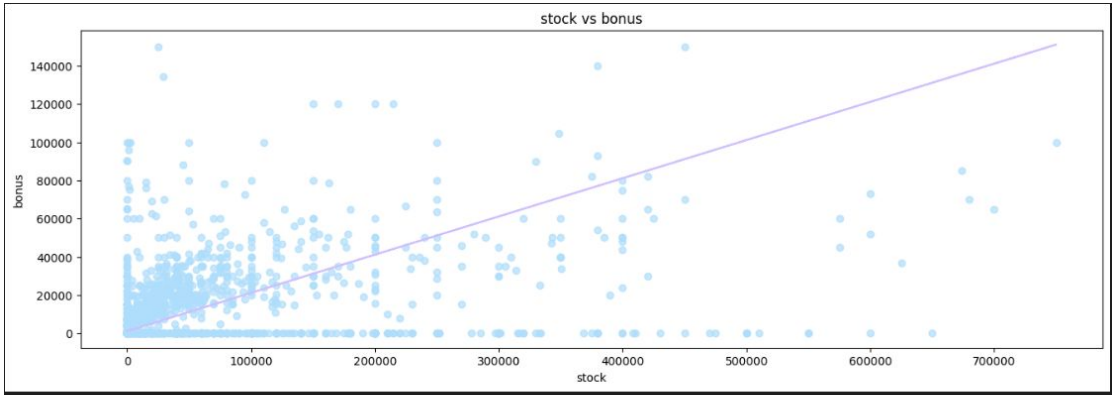

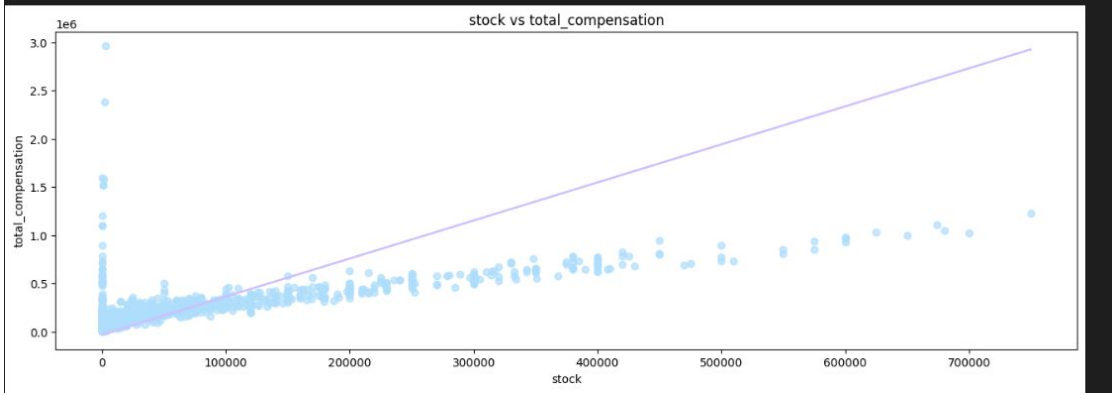

Visualization

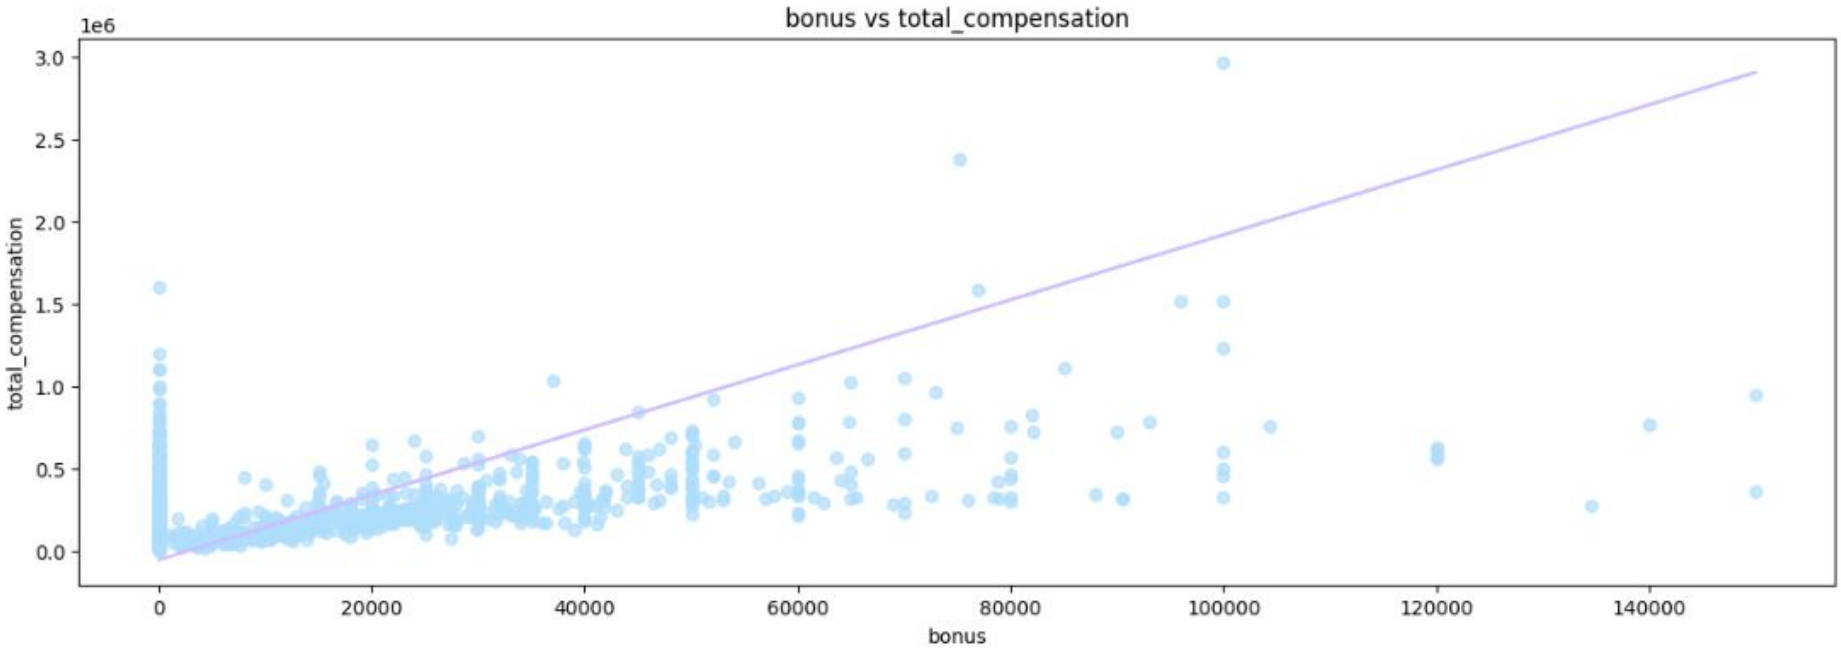

Data Modeling

Define problem

Define what problem to be solved

Prepare Data

Prepare data for training and testing

Visualization

Visualize the result

Feature Engineering

Encode categorical features

Modeling

Create model, train and test model

Define problem

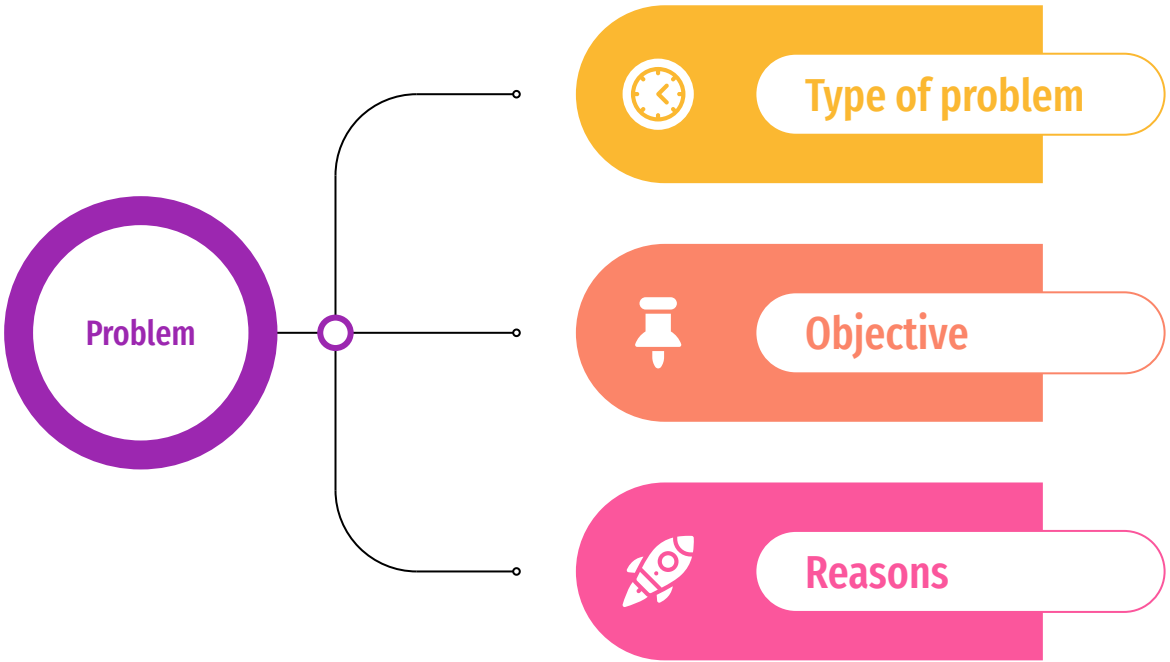

Regression

The total compensation of an employee based on various job-related features

predict a continuous value, specifically the total compensation of an employee

Data Preparing

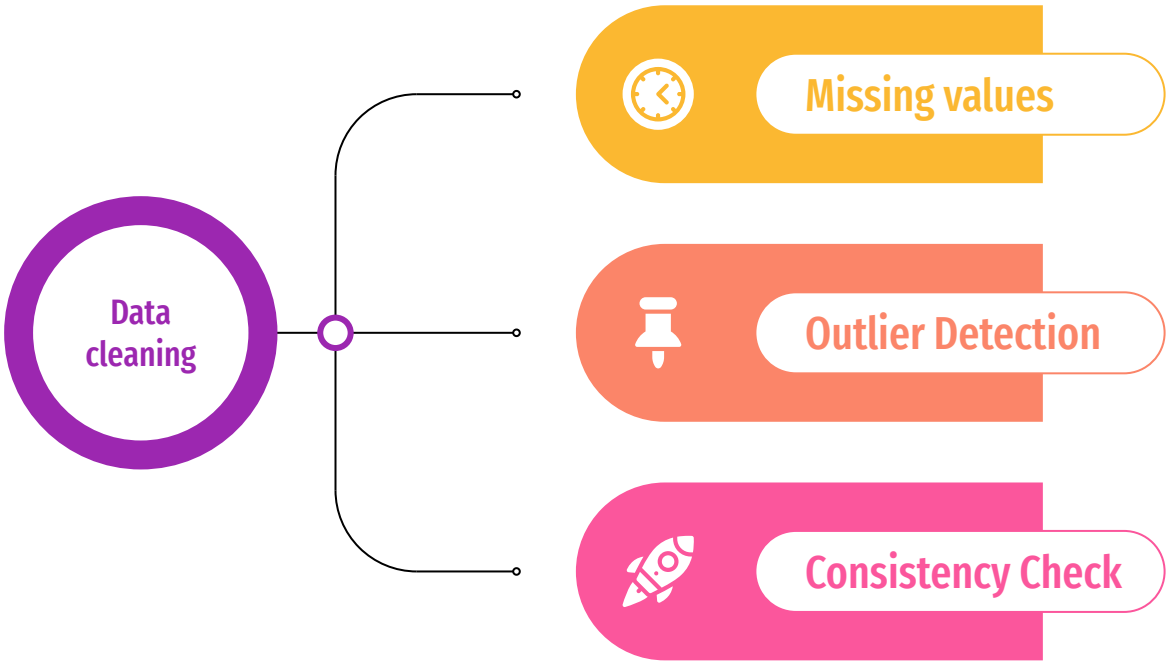

Replace them with the mean or mode

Calculate z-score and replace the outlier with the mean

Set all negative values to zero

Feature Engineering

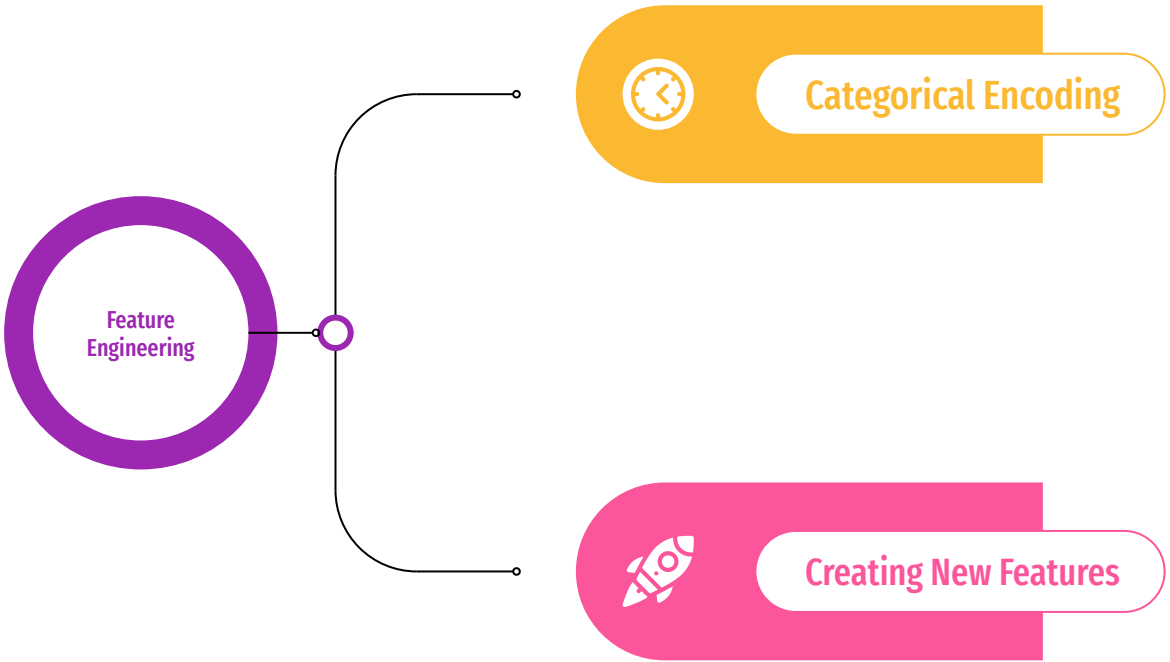

Encode categorical columns

Provide more information to train and improve the performance of the model.

Modeling

01

Prepare training and testing data

02

Training data with models

03

Evaluate the results

04

Visualize the results

Modeling

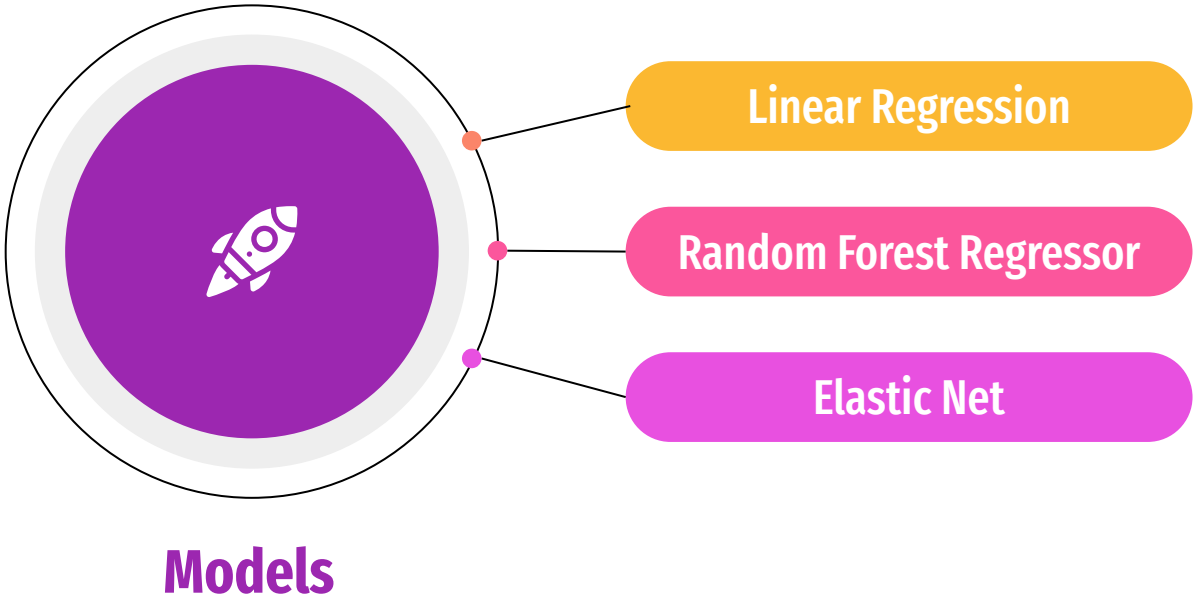
Linear Regression
Random Forest Regressor
Elastic Net
Models

Modeling

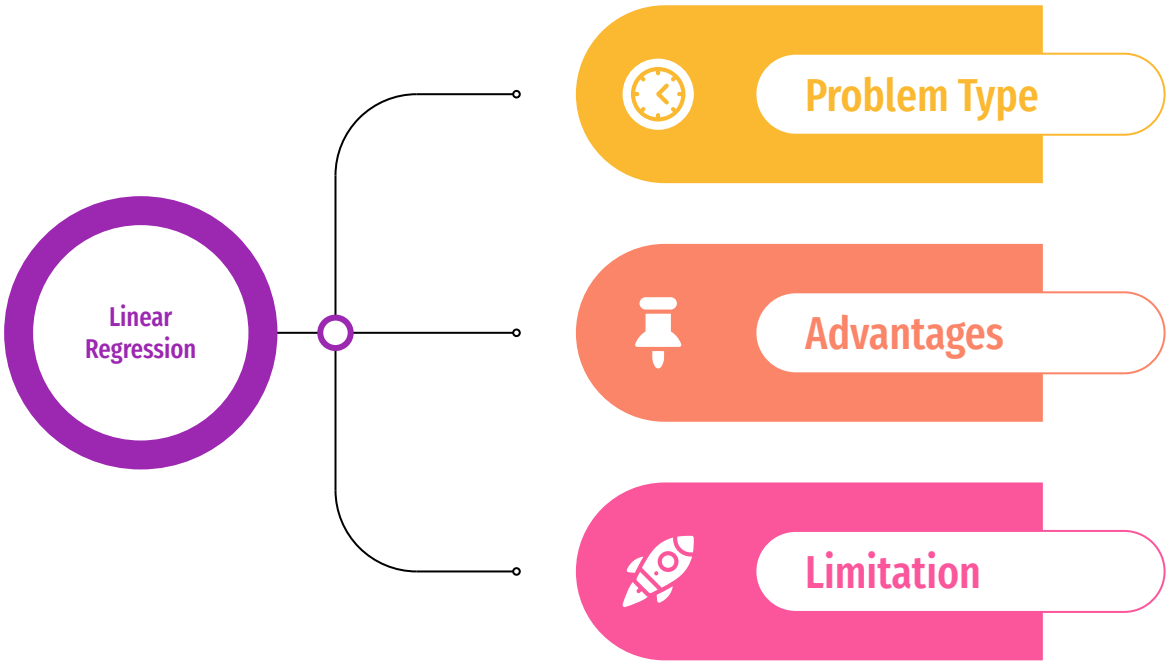
Linear Regression
Problem Type
linear relationship between the (in)dependent variables
Advantages
Simple, easy, and effective when the relationship is linear
Limitation
handle non-linear relationships and complex

Modeling

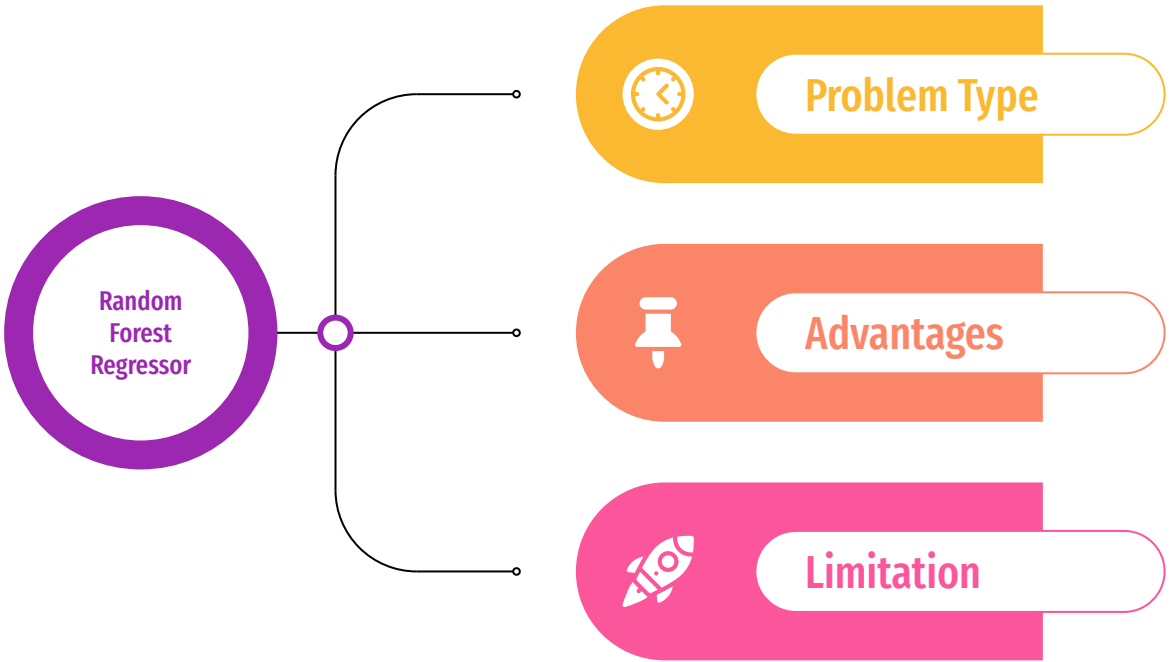

both regression and classification

Flexible, less sensitive to noise and overfitting.

Requires a sufficiently large dataset to be effective.

Modeling

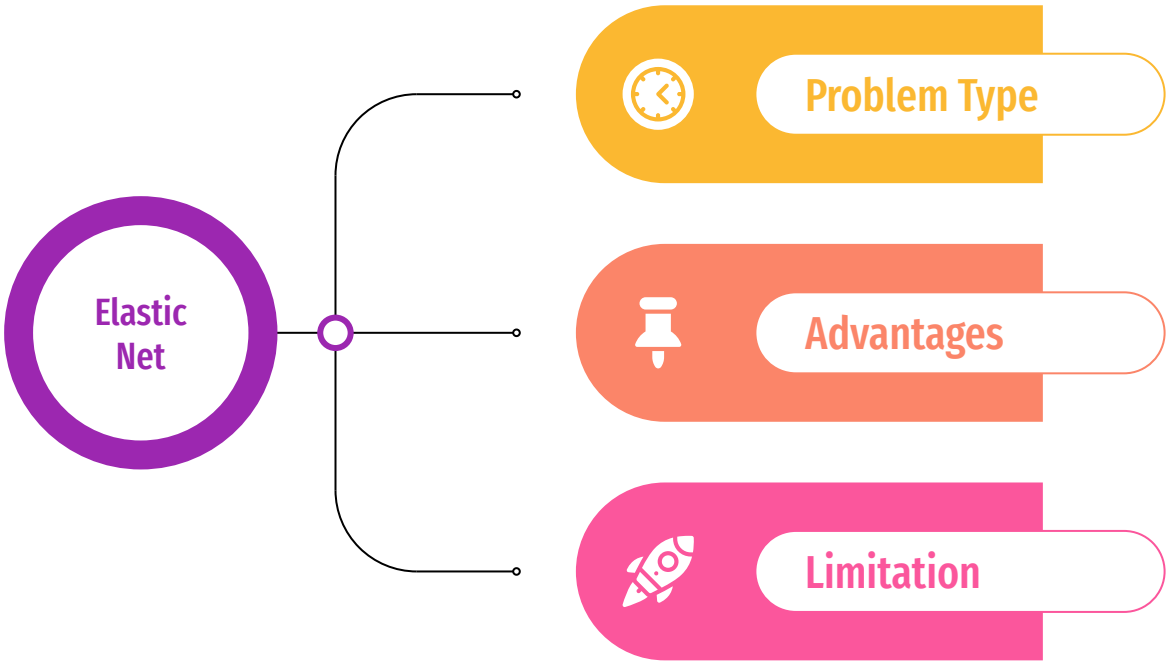
Elastic Net
Problem Type
linear relationship and high correlation between features
Advantages
reducing overfitting and retaining important features
Limitation
May require tuning of hyperparameters for optimal performance

Compare the results

Linear Regression

```
Mean Squared Error: 3852930.7381835002
R-squared: 0.9998159615414253
```

Random Forest

```
Mean Squared Error: 134939215.97234216
R-squared: 0.9935545155113429
```

Elastic Net

```
Mean Squared Error: 3837235.8017152497
R-squared: 0.9998167112231899
```

Evaluate and compare

Conclusion

-> After fine-tuning and re-training the models on the combined training and validation sets, the performance metrics, particularly Mean Squared Error (MSE) and R-squared, demonstrated notable improvement. This suggests that the models have effectively learned from a larger and more diverse dataset, resulting in enhanced predictive capabilities on new, unseen data.

-> Based on the information from the MSE and R-squared values of all three models, it is evident that the Elastic Net model demonstrates significant effectiveness, as it exhibits the lowest MSE and the highest R-squared among the models.

Visualization

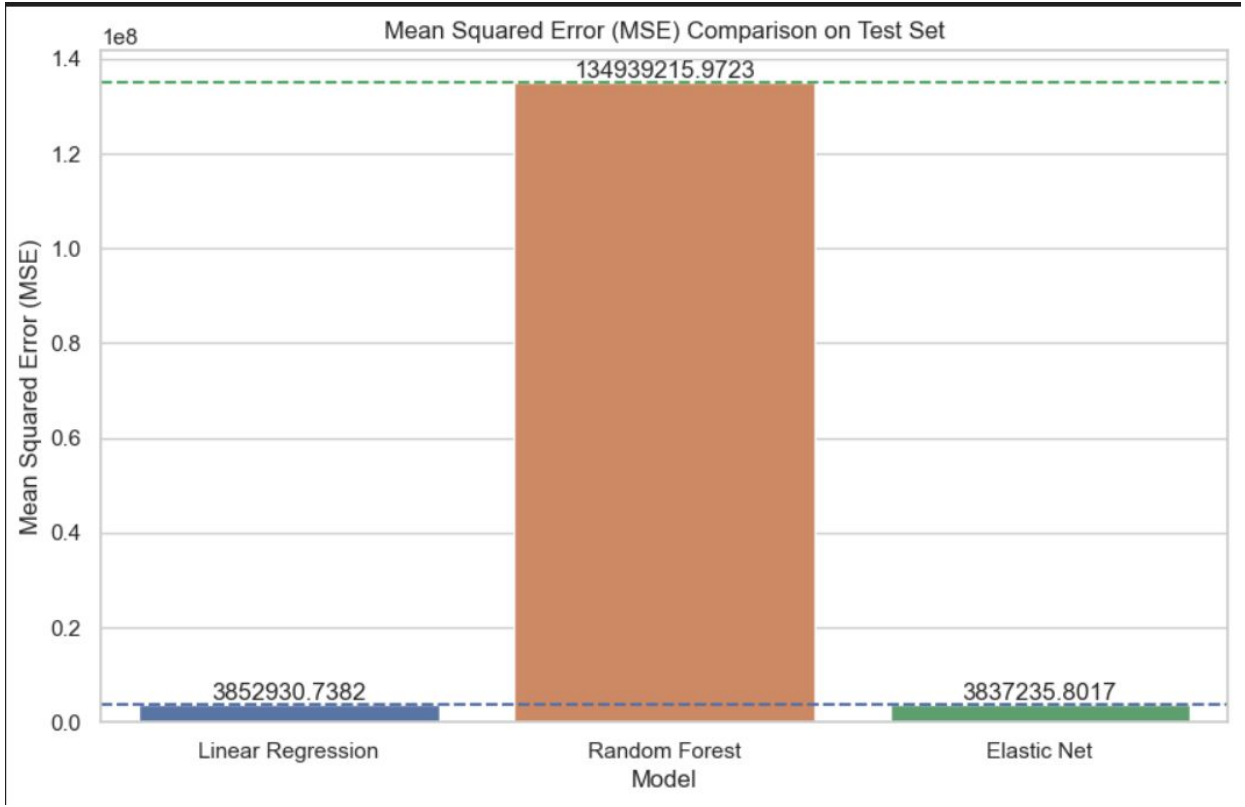

Visualization

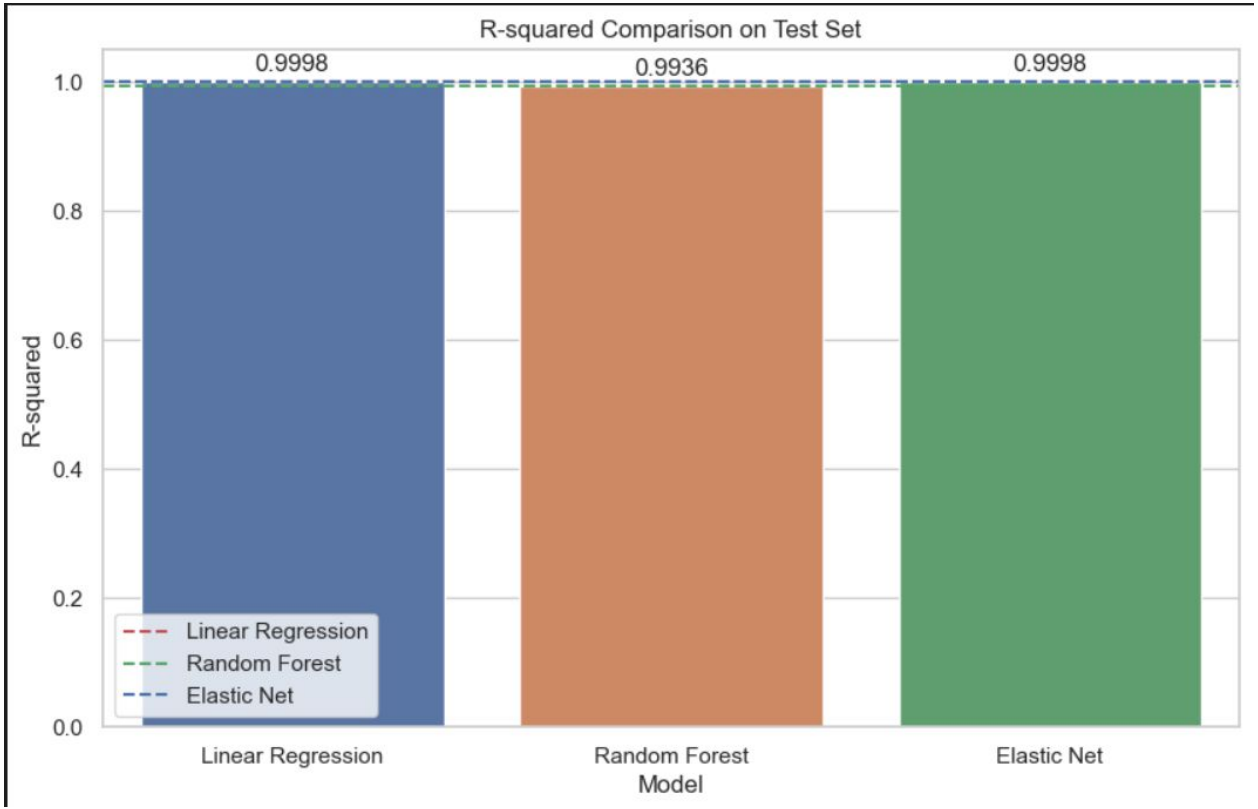

Thanks for your listening!